

SESSION 7 · MINI-ASSESSMENT REPORT

HappySoul – User Emotional Check-In Triage

Design and Implementation of a Multi-Step Prompt Workflow

Name	GUTTULA JASWANTH
Student ID	99230040110
Programme	Gen AI
Module	Prompt Engineering
Session	Session 7 – Multi-Step Prompt Workflow
Date of Submission	05-09-2026

Table of Contents

1. Introduction..... 3

2. Objective..... 3

3. Workflow Design..... 3

4. Methodology..... 4

 4.1 Step 1: Emotional Entity Extraction Using Few-Shot Prompting..... 4

 4.2 Step 2: Response-Tier Assessment Using Chain-of-Thought Reasoning 5

 4.3 Step 3: Structured Output Using JSON..... 6

5. Results and Discussion 7

6. Conclusion 7

1. Introduction

This report presents the design and implementation of a three-step prompt workflow for HappySoul, a spiritual wellness app whose Krishna AI companion responds to users' mood and journal check-ins. The workflow takes an unstructured check-in written by a user and progressively transforms it into a structured result that HappySoul's chatbot layer can act on.

The workflow follows this sequence:

Raw Check-In → Emotional Entity Extraction → Response-Tier Assessment → Structured JSON Output

Each stage applies a prompting technique suited to its purpose:

- Few-shot prompting to extract the emotional details that matter from the check-in.
- Chain-of-thought reasoning to decide how deep and careful the AI companion's response should be.
- JSON formatting to produce a structured, machine-readable output the chatbot backend can consume.

2. Objective

- Extract the emotionally relevant information from an unstructured check-in using few-shot prompting.
- Analyse the extracted information and determine the appropriate response tier using structured reasoning.
- Convert the results into a valid JSON object that Krishna AI's response layer could process.
- Demonstrate that the output of each stage is used as the input for the next.

3. Workflow Design




Figure 1: Prompt Chaining Workflow for HappySoul's Check-In Triage

Separating the workflow into three stages ensures that each prompt has a single, focused responsibility. Step 1 identifies what the user is actually feeling and why; Step 2 makes a judgement about how carefully the reply needs to be handled; and Step 3 packages the result for the chatbot backend.

This separation keeps each individual prompt simple and makes the workflow easier to inspect — an important property for a wellness app, since every intermediate judgement can be reviewed on its own rather than trusting one opaque response.

4. Methodology

The workflow was implemented using Google AI Studio. The following user check-in was used as the input for the experiment:

“I lost my job today and I still haven't figured out how to tell my family. I keep replaying the meeting in my head and I can't stop crying. I don't really want to talk to anyone about it right now, I just needed to write it down somewhere.”

4.1 Step 1: Emotional Entity Extraction Using Few-Shot Prompting

Purpose

The purpose of the first stage was to extract the most important emotional information from the user's unstructured check-in. Few-shot prompting was chosen because worked examples demonstrate the relationship between an input and its desired output, while establishing a consistent format for the extracted information. Three examples spanning a calm entry, a moderately distressed entry, and a high-concern entry were provided before the actual check-in.

Prompt Used

I'm building the entity-extraction stage of an emotional check-in triage workflow for a spiritual wellness app called HappySoul. For this step, I need you to read a user's mood/journal check-in and pull out the details that matter for deciding how the app's Krishna AI companion should respond.

Use the examples below as a guide for what to look for and how to format your answer.

Example 1:

User check-in: “Today was a decent day, nothing special happened, but I felt really calm during my evening walk.”

Output:

- Emotional state: calm, content
- Trigger/context: evening walk, routine day
- Intensity: low
- Crisis language: none
- User tone: reflective, at ease

Example 2:

User check-in: “I've been so anxious about my exams. I can't focus on anything and I feel like I'm letting everyone around me down.”

Output:

- Emotional state: anxiety, self-doubt
- Trigger/context: upcoming exams, academic pressure
- Intensity: moderate to high
- Crisis language: none
- User tone: overwhelmed, self-critical

Now analyse this check-in using the same style:

“I lost my job today and I still haven't figured out how to tell my family. I keep replaying the meeting in my head and I can't stop crying. I don't really want to talk to anyone about it right now, I just needed to write it down somewhere.”

Output Generated

- Emotional state: shock, grief, shame
- Trigger/context: sudden job loss, anticipated difficult conversation with family
- Intensity: high
- Crisis language: none (no expressions of hopelessness or self-harm)
- Previous coping: chose to write it down rather than talk to someone
- User tone: distressed, currently withdrawing from others

4.2 Step 2: Response-Tier Assessment Using Chain-of-Thought Reasoning

Purpose

The purpose of the second stage was to determine how carefully Krishna AI should handle its reply. Unlike entity extraction, judging response tier requires weighing several factors together, so the workflow used a structured reasoning approach that considers safety risk, how much the distress is affecting the user right now, its intensity, and whether the user has already sought support elsewhere.

Prompt Used

Now that the important details have been extracted, the next step is to decide how HappySoul's Krishna AI should respond to this user in terms of depth and tone.

Look at the information below and reason through the decision carefully. Consider whether there is any crisis language indicating risk to the user's safety, how much the issue is affecting the user's ability to cope right now, how intense and recent the distress is, and whether the user has already tried to reach out to someone else.

Use the following guide:

- Reflective / Routine: the entry is calm or neutral, with no significant distress; a light, encouraging response is enough.
- Supportive Guidance: a genuine but manageable difficulty is present; the user would benefit from an encouraging message and a relevant Gita theme.
- Elevated Emotional Support: the distress is significant and affecting the user's day, but there is no indication of risk to safety; the response should be deeper, validating, and paired with a specific Gita verse and a gentle follow-up prompt.
- Crisis / Safety Escalation: the entry contains language suggesting risk to the user's safety or wellbeing; the app must pause the normal guidance flow and surface professional crisis-support resources instead.

Explain your decision in a few numbered points. Base your answer only on the information provided below.

Extracted check-in details:

- Emotional state: shock, grief, shame
- Trigger/context: sudden job loss, anticipated difficult conversation with family
- Intensity: high
- Crisis language: none
- Previous coping: chose to write it down rather than talk to someone

- User tone: distressed, currently withdrawing from others

Output Generated

1. The entry shows acute distress (shock, grief, shame) triggered by a specific, significant life event, rather than a general or ongoing mood dip.
2. There is no language suggesting hopelessness, self-harm, or a wish to give up, so this does not meet the criteria for crisis escalation.
3. The user is currently avoiding talking to others about it, which suggests they may not have another emotional outlet in this moment, so a brief or generic response would likely feel dismissive.

Response tier: Elevated Emotional Support

Discussion

The check-in was classified as Elevated Emotional Support because the distress was acute and clearly tied to a specific event, no safety-risk language was present, and the user's current avoidance of talking to others suggested that a shallow response would land poorly.

This reasoning stage worked from the information extracted in Step 1 rather than re-reading the original check-in independently, demonstrating the chaining relationship between the two prompts. The final tier decision was then passed to the third stage.

4.3 Step 3: Structured Output Using JSON

Purpose

The purpose of the final stage was to convert the results of the previous two stages into a structured format. JSON was chosen because it is consistently structured and easy for software systems to process — a result such as this is what HappySoul's Krishna AI response layer would consume to decide what to say and which Gita content to surface.

Prompt Used

Convert the results from the previous two stages of this HappySoul emotional check-in triage workflow into a valid JSON object.

Return only the JSON object. Do not include explanations, headings, or markdown.

Use exactly these three keys:

- “entities”: a list of the important details extracted from the check-in
- “response_tier”: the final response tier decided during the previous step
- “recommended_action”: one to two sentences describing what Krishna AI should say or do, and why that fits the tier

Results from Step 1:

- Emotional state: shock, grief, shame
- Trigger/context: sudden job loss, anticipated difficult conversation with family
- Intensity: high
- Crisis language: none
- Previous coping: chose to write it down rather than talk to someone

- User tone: distressed, currently withdrawing from others

Result from Step 2:

Response tier: Elevated Emotional Support

Final JSON Output Generated

```
{
  "entities": [
    "Emotional state: shock, grief, shame",
    "Trigger/context: sudden job loss, anticipated difficult conversation with family",
    "Intensity: high",
    "Crisis language: none",
    "Previous coping: chose to write it down rather than talk to someone",
    "User tone: distressed, currently withdrawing from others"
  ],
  "response_tier": "elevated_emotional_support",
  "recommended_action": "Krishna AI should respond with a validating message acknowledging the shock and shame of losing a job, gently offer a Gita verse on detachment from outcomes and self-worth beyond one's role (e.g. Bhagavad Gita 2.47), and softly suggest logging today's mood for later, without pressuring the user to talk to anyone before they are ready."
}
```

5. Results and Discussion

The three-stage workflow successfully processed the original check-in, carrying it from an unstructured journal entry through to a structured JSON result that HappySoul's backend could act on directly.

Table 1: Summary of the Prompt Workflow

Step	Task	Prompting Technique	Output
Step 1	Extract relevant emotional details	Few-shot prompting	Structured emotional-entity list
Step 2	Determine response tier	Chain-of-thought reasoning	Response-tier classification
Step 3	Format the results	JSON output	Machine-readable JSON object

This separation makes the workflow easier to review, since each intermediate result can be inspected on its own — which matters for a wellness product, where it should always be possible to see why the AI chose a particular tone. It also means the final output follows a consistent schema containing the three required fields:

- entities
- response_tier
- recommended_action

The resulting JSON object is well-formed and can be parsed by a standard JSON parser without requiring manual cleanup. Because the tier classification is grounded in explicit, inspectable reasoning (Step 2), the workflow could also be extended so

that a Crisis / Safety Escalation tier automatically routes the user to professional support resources rather than letting the chatbot continue the conversation on its own.

6. Conclusion

This approach offers greater transparency, since intermediate outputs can be reviewed individually, and illustrates how large language models could be integrated into a wellness product like HappySoul to make Krishna AI's responses both emotionally appropriate and safe to escalate when needed.

- Few-shot prompting was used to extract emotionally relevant information from an unstructured user check-in.
- Chain-of-thought reasoning was used to weigh safety-risk language, current coping, intensity, and prior help-seeking to determine the appropriate response tier.
- JSON formatting was used to convert the final results into a structured, machine-readable output.